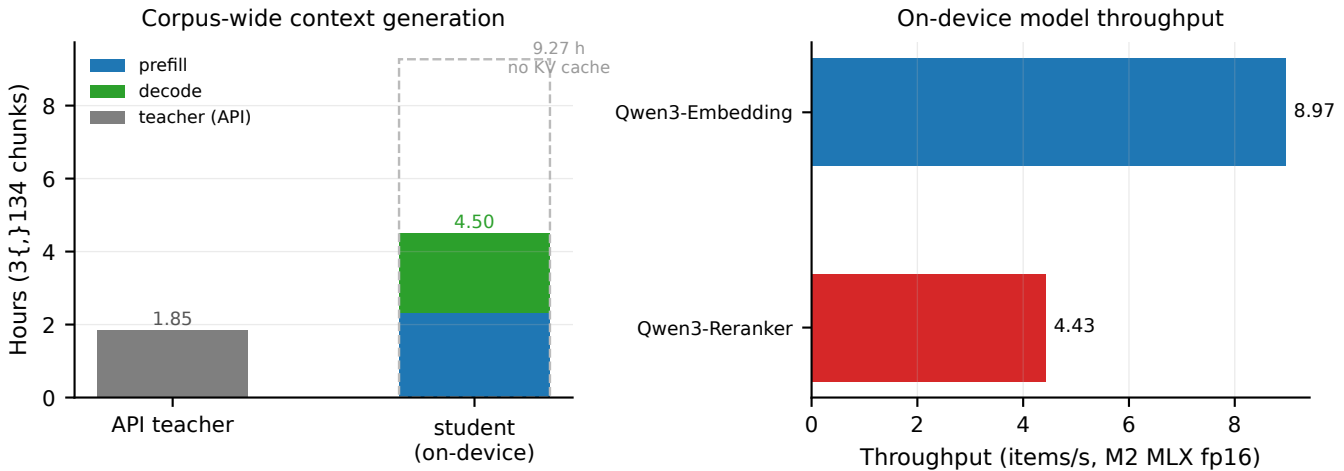
Corpus-wide context generation
prefill
decode
teacher (API)
9.27 h
no KV cache
4.50
1.85
Hours (3{,}134 chunks)
0
2
4
6
8
API teacher
student
(on-device)
On-device model throughput
Qwen3-Embedding
8.97
Qwen3-Reranker
4.43
0
2
4
6
8
Throughput (items/s, M2 MLX fp16)